

Machine learning with Galaxy


Qiang Gu and Anup Kumar

European Galaxy Days, 2018
Faculty of Engineering,
Freiburg

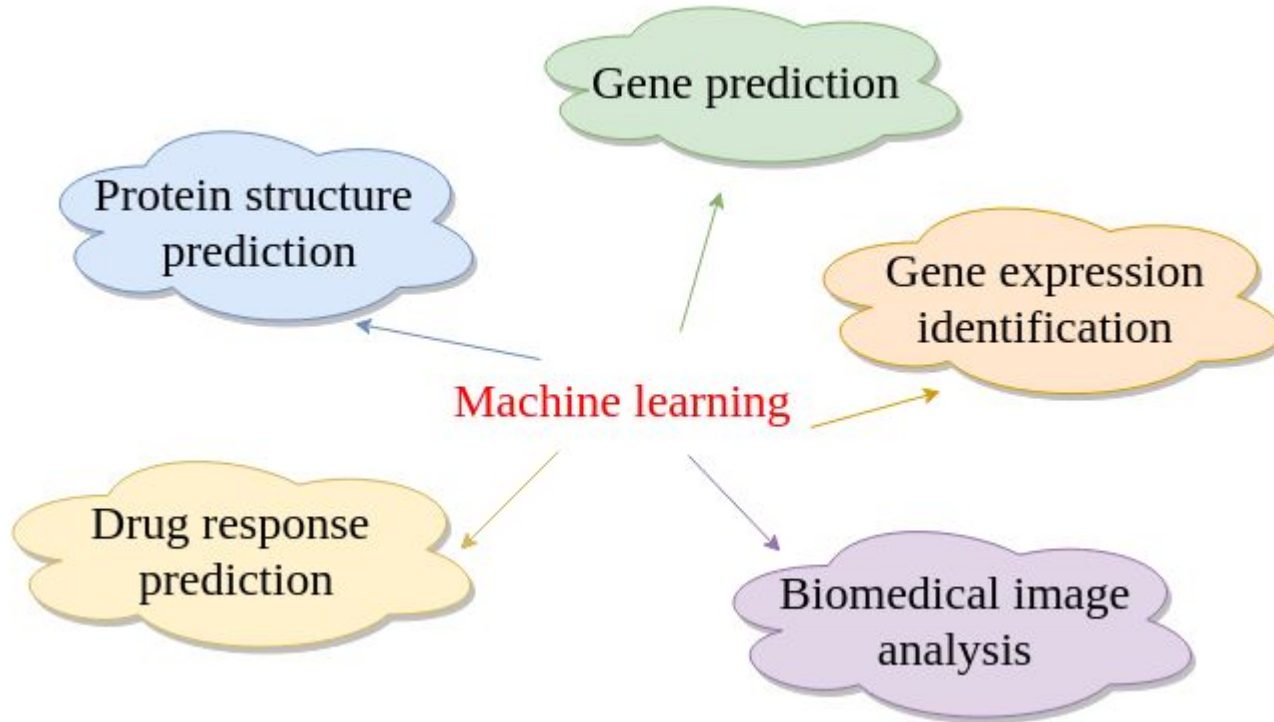
Outline

- Machine learning
- Significance in Bioinformatics
- Galaxy's machine learning tools
- Use-case: Regression
 - Penn Machine Learning Benchmark datasets

Machine learning



Application in Bioinformatics



Scikit-learn in Galaxy

- Scikit-learn modules as Galaxy tools
- Scikit-learn:
 - Open source, python
 - 89 releases
 - 1,200+ contributors
 - 15,000+ forks
 - 30,000+ stargazers



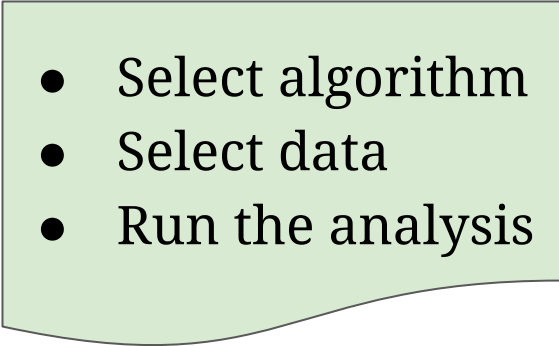
Various machine learning tools in Galaxy

- 20+ classifiers (classification)
- 20+ regressors (regression)
- 30+ data preprocessing techniques
- 2 hyperparameter search approaches
 - Grid search
 - Random search
- 2+ plotting tools
 - ROC, AUC, confusion matrix
 - True vs predicted, residual plot

- Linear models
- Support vector machines
- Nearest neighbours
- Tree and ensemble methods
- Naive bayes
- ...

Usability and accessibility

- No programming knowledge required
- Easy to use UI
- Create workflows for end-to-end analysis
- Parallel processing (a few modules)
- Submit jobs (dataset collection)
- Save and reuse trained models
- 1 training material online, 1 in progress

- 
- Select algorithm
 - Select data
 - Run the analysis

Tool

Select an ensemble method:

Gradient Boosting Regressor

Select input type:

tabular data

Training samples dataset:



No tabular dataset available.

Does the dataset contain header:

Yes

No

Choose how to select data by column:

Select columns by column index number(s)

Select target column(s):

Dataset containing class labels or target values:



No tabular dataset available.

Does the dataset contain header:

Yes

No

Choose how to select data by column:

Select columns by column index number(s)

Select target column(s):

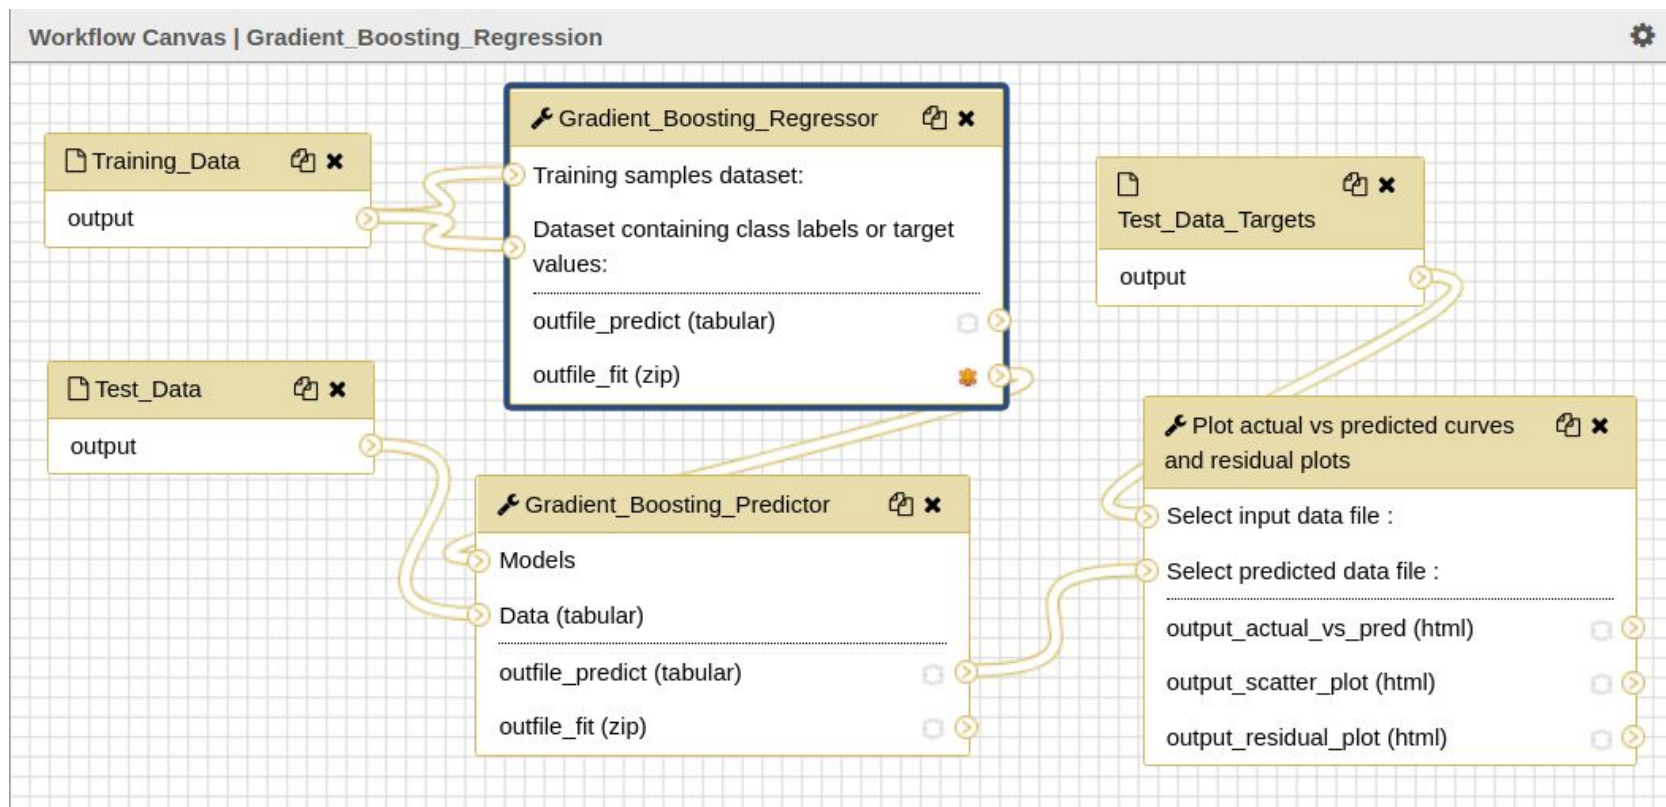
Select algorithm

Select data

Done!

8

Workflow



Hyperparameters

Advanced Options

Loss function

ls - least squares regression

(loss)

Learning rate

0.1

(learning_rate)

Number of trees in the forest

100

The number of boosting stages to perform (n_estimators)

Maximum depth of the tree

3

maximum depth of the individual regression estimators (max_depth)

Function to measure the quality of a split

friedman_mse - mean squared error with improvement score by Friedman

(criterion)

Minimum number of samples required to split an internal node

2.0

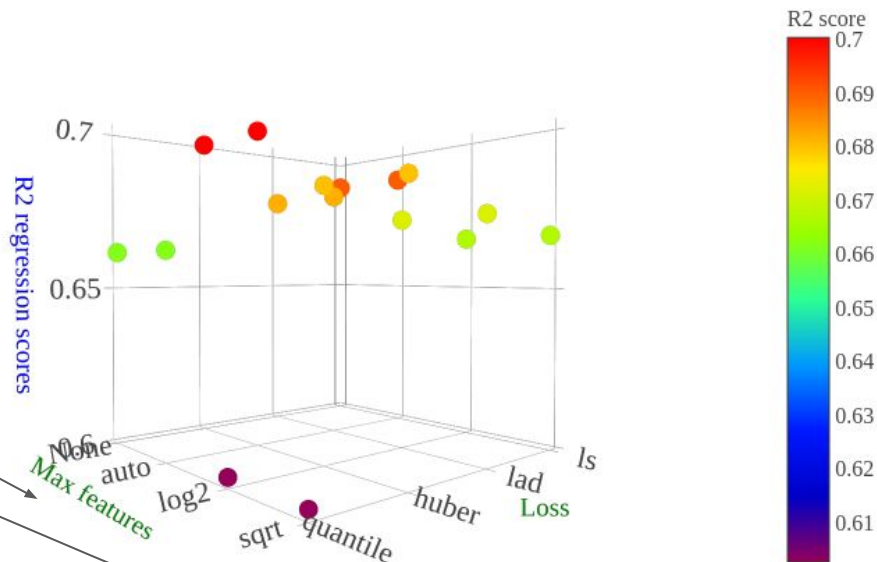
- Loss function
- Learning rate
- Number of trees
- Maximum depth of tree
- Function to measure quality of split
- Minimum samples to split

Hyperparameter optimisation

- Grid search

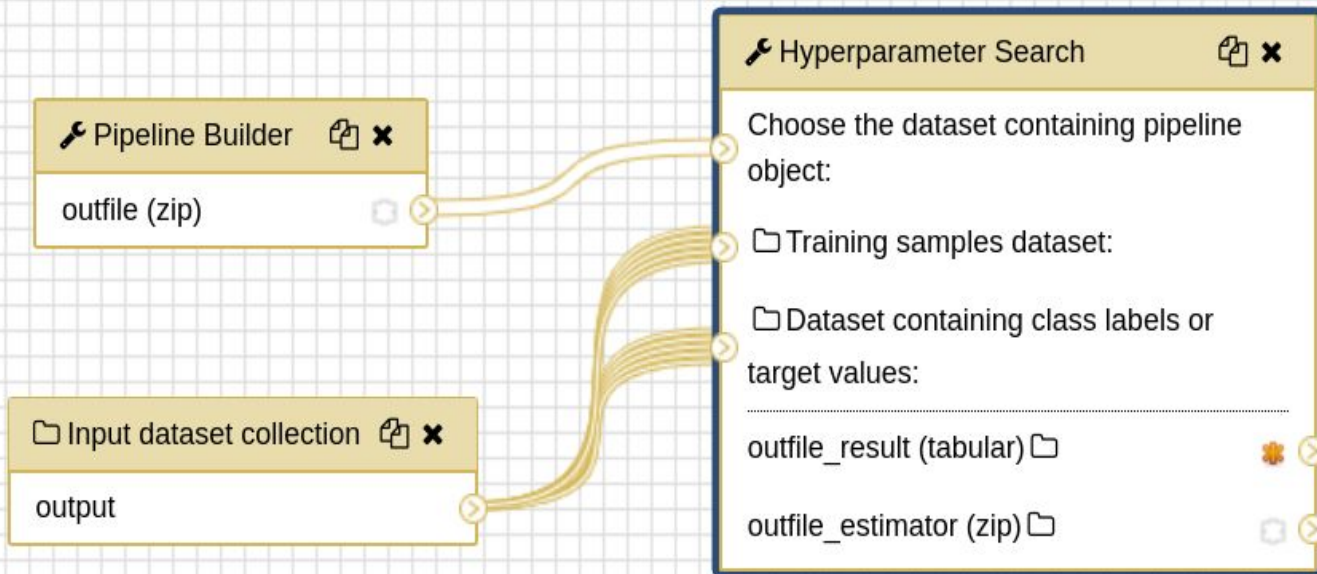
models: $4 \times 4 = 16$

Gradient Boosting Regression



Hyperparameter optimisation

Workflow Canvas | GradientBoosting_Regressor



1. Create pipeline

Workflow Canvas | GradientBoosting_Talk_Regressor

Hyperparameter Search

Choose the dataset containing pipeline object:

- Training samples dataset:
- Dataset containing class labels or target values:

outfile_result (tabular)

outfile_estimator (zip)

Details

Pre-processing step

1: Pre-processing step

Choose the type of transformation:

Feature Selection

Select a feature selection algorithm

SelectKBest - Select features according to the...

Select a score function

f_regression - Univariate linear regression t...
(score_func)

Advanced Options

Final Estimator

Choose the module that contains target estimator:

sklearn.ensemble

Choose estimator class:

GradientBoostingRegressor

- Select preprocessing
- Select algorithm

2. Set hyperparameters

2: Parameter setting for search:

Choose the transformation the parameter belongs to

Final estimator

☒ Estimator parameter:

loss: ['ls', 'lad', 'huber', 'quantile']

One parameter per box. For example: C: [1, 10, 100, 1000]. See bottom for more examples

- 4 types of losses
- 4 types of max features

3: Parameter setting for search:



Choose the transformation the parameter belongs to

Final estimator

☒ Estimator parameter:

max_features: [None, 'auto', 'log2', 'sqrt']

One parameter per box. For example: C: [1, 10, 100, 1000]. See bottom for more examples

3. Search results

Accuracy

Hyperparameters

Models

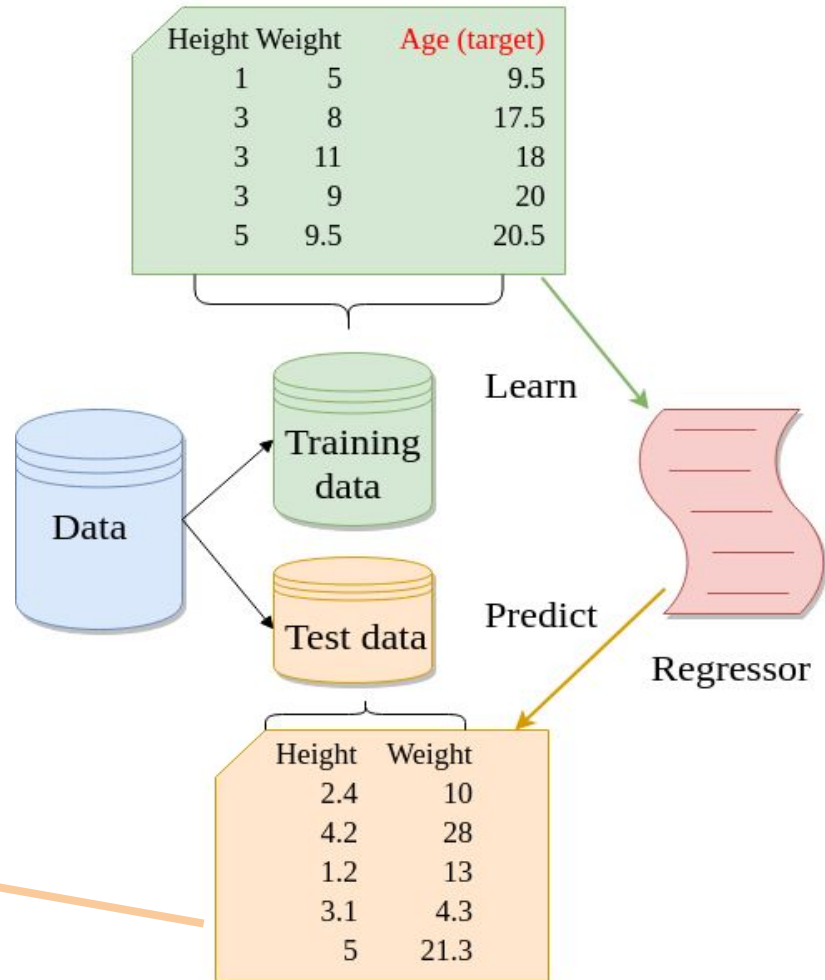
3	4	5	6	7
mean_test_score	param_estimator__loss	param_estimator__max_features	param_estimator__random_state	params
0.6903816799925716	ls		3111696	{'estimator__loss': 'ls', 'estimator__max_features': None, 'estimator__random_state': 3111696}
0.6903816799925716	ls	auto	3111696	{'estimator__loss': 'ls', 'estimator__max_features': 'auto', 'estimator__random_state': 3111696}
0.6667323248740175	ls	log2	3111696	{'estimator__loss': 'ls', 'estimator__max_features': 'log2', 'estimator__random_state': 3111696}
0.6667323248740175	ls	sqrt	3111696	{'estimator__loss': 'ls', 'estimator__max_features': 'sqrt', 'estimator__random_state': 3111696}
0.6817019314143619	lad		3111696	{'estimator__loss': 'lad', 'estimator__max_features': None, 'estimator__random_state': 3111696}
0.6817019314143619	lad	auto	3111696	{'estimator__loss': 'lad', 'estimator__max_features': 'auto', 'estimator__random_state': 3111696}
0.6716985021609994	lad	log2	3111696	{'estimator__loss': 'lad', 'estimator__max_features': 'log2', 'estimator__random_state': 3111696}
0.6716985021609994	lad	sqrt	3111696	{'estimator__loss': 'lad', 'estimator__max_features': 'sqrt', 'estimator__random_state': 3111696}
0.7003788033314727	huber		3111696	{'estimator__loss': 'huber', 'estimator__max_features': None, 'estimator__random_state': 3111696}
0.7003788033314727	huber	auto	3111696	{'estimator__loss': 'huber', 'estimator__max_features': 'auto', 'estimator__random_state': 3111696}
0.6800040571751129	huber	log2	3111696	{'estimator__loss': 'huber', 'estimator__max_features': 'log2', 'estimator__random_state': 3111696}
0.6800040571751129	huber	sqrt	3111696	{'estimator__loss': 'huber', 'estimator__max_features': 'sqrt', 'estimator__random_state': 3111696}
0.6616117433548167	quantile		3111696	{'estimator__loss': 'quantile', 'estimator__max_features': None, 'estimator__random_state': 3111696}
0.6616117433548167	quantile	auto	3111696	{'estimator__loss': 'quantile', 'estimator__max_features': 'auto', 'estimator__random_state': 3111696}
0.6024828194032356	quantile	log2	3111696	{'estimator__loss': 'quantile', 'estimator__max_features': 'log2', 'estimator__random_state': 3111696}
0.6024828194032356	quantile	sqrt	3111696	{'estimator__loss': 'quantile', 'estimator__max_features': 'sqrt', 'estimator__random_state': 3111696}

**Regression
on
Penn Machine Learning Benchmark datasets**

Regression


- Supervised learning
- Real valued targets (output)
- R2 scoring metric

Predict “Age”




Regression metric (R2)

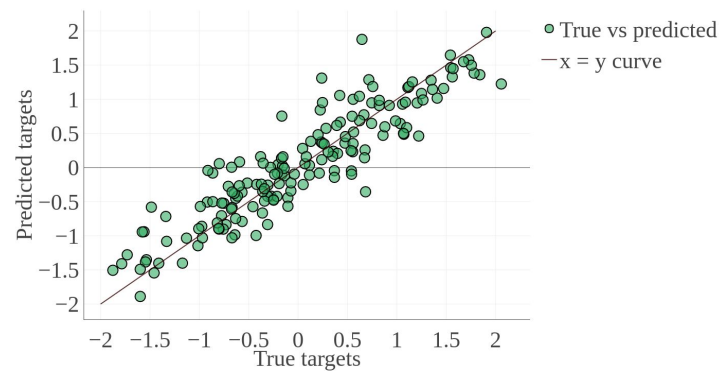
True vs predicted targets (R2: -0.33)




True vs predicted targets (R2: 0.44)




True vs predicted targets (R2: 0.82)




True vs predicted targets (R2: 1.00)




R2 regression scores vs regressors



R2 scores of regressors with and without hyperparameter optimisation



Fit time vs R2 regression score



Conclusion

- Machine learning with Galaxy
- Variety of preprocessing and learning techniques
- Applications in Bioinformatics
- Usage:
 - Create workflows
 - Optimise hyperparameters
 - Analyse results using plots
 - Regression on a public dataset

Thank you for your attention

Questions?

References

- Original paper [<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC5890912/>]
- Machine learning in Bioinformatics [<https://academic.oup.com/bib/article/7/1/86/264025>]
- Drug prediction [<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4396063/>]
- Penn Machine Learning Benchmarks [<https://github.com/EpistasisLab/penn-ml-benchmarks>]
- Protein structure prediction with machine learning
[<https://www.ncbi.nlm.nih.gov/pubmed/22274898>]
- Gene prediction [<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC5698827/>]
- Machine learning and genome annotation
[<https://genomebiology.biomedcentral.com/articles/10.1186/gb-2013-14-5-205>]
- Machine learning on gene expression microarray data
[<https://www.ncbi.nlm.nih.gov/pubmed/18366602>]
- Medical image analysis with machine learning
[<https://www.sciencedirect.com/science/article/pii/S1361841516301098>]
- First training material
[<https://galaxyproject.github.io/training-material/topics/statistics/tutorials/machinelearning/tutorial.html>]